

Lecture 10: Deep Sequence Modeling. Recurrent neural networks.

Predictive learning

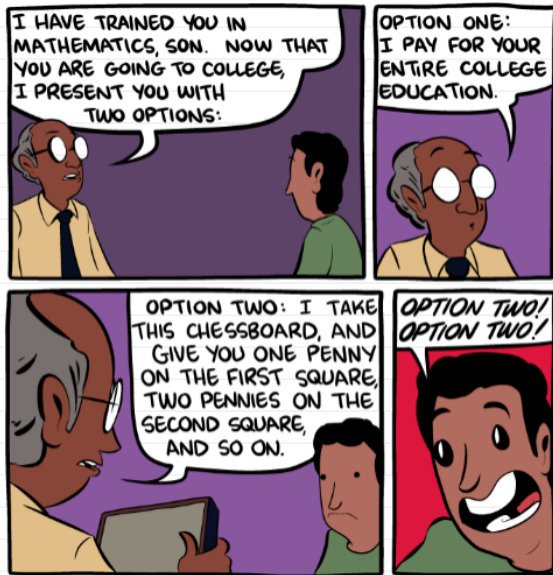
- Given an element predict *nearby* elements (e.g. next, previous, adjacent, etc.)
- Does not require annotated data (“self-supervised”)
- Usually considered as unsupervised, but often works much better than “plain” unsupervised
- Particularly prominent in NLP, but now gaining popularity in many fields

Today's focus: sequence modeling,
sequence prediction

Predicting sequences matters

Applications:

- Synthesis (text, speech, etc.)
- Probabilistic modelling
- Compression



smbc-comics.com

Training sequence prediction

A cat sat on a ma?



Inherently probabilistic: need to predict probabilities over alphabet/lexicon

Training sequence prediction

A cat sat on a ma?

Predominantly maximum likelihood learning:

$$\max_{\theta} \sum_i \log p_{\theta}(x_t^i | x_{t-1}^i, x_{t-2}^i, \dots, x_1^i)$$

Many models go back fixed number of steps:

$$\max_{\theta} \sum_i \log p_{\theta}(x_t^i | \underbrace{x_{t-1}^i, x_{t-2}^i, \dots, x_{t-N}^i}_{\text{Temporal window}})$$

Fixed window/order architectures

$$p_{\theta}(x_t^i \mid x_{t-1}^i, x_{t-2}^i, \dots, x_{t-N}^i)$$

- N-grams (with smoothing)
- ConvNets (aka TDNNs)
- Any probabilistic classifier (e.g. decision forest, etc.)

NB: using padding for the special symbol (UNK) we can train model for shorter sequences

Assessing a probabilistic model

1. Train $p_{\theta}(x_j \mid x_{j-1}, \dots, x_{j-N})$

2. Evaluate $\prod_{j=1}^M p_{\theta}(x_j \mid x_{j-1}, \dots, x_{j-N})$
on a hold-out set (can be a long text)

Common measure (*perplexity*):

$$\text{PP}(x_1, \dots, x_M) = \sqrt[M]{\prod_{j=1}^M p_{\theta}(x_j \mid x_{j-1}, \dots, x_{j-N})}^{-1}$$

Probabilistic modeling of long sequences

Assume given $p_{\theta}(x_t^i \mid x_{t-1}^i, x_{t-2}^i, \dots, x_{t-N}^i)$

$$p(x_M, x_{M-1}, \dots, x_1) =$$

$$p(x_M \mid x_{M-1}, \dots, x_1) \cdot p(x_{M-1}, x_{M-2}, \dots, x_1) =$$

$$\prod_{j=2}^M p(x_j \mid x_{j-1}, \dots, x_1) \cdot p(x_1) \approx$$

$$\prod_{j=1}^M p_{\theta}(x_j \mid x_{j-1}, \dots, x_{j-N})$$

ML sequence generation

Task: draw a sample sequence with high-probability


$$\prod_{j=1}^M p_{\theta}(x_j \mid x_{j-1}, \dots, x_{j-N})$$

Option 1: synthesize one-by-one greedily, picking the symbol with highest probability

$$\hat{x}_j = \arg \max_x p_{\theta}(x \mid \hat{x}_{j-1}, \dots, \hat{x}_{j-N})$$

Option 2: *beam search*

Why greedy synthesis is suboptimal



Toy example: three letters in the alphabet.
Task: synthesize most likely three letter word starting from red.

Greedy solution:



Best solution:



Beam search

The c??????	$\prod_{j=1}^M p_{\theta}(x_j \mid x_{j-1}, \dots, x_{j-N})$	
The ca?????	<u>The cat????</u> <u>The cap????</u>	The cat????
The co?????	<u>The cor????</u> The col????	The cap????
The ch?????	The cha???? The cho????	The cor????


WaveNet



1 Second


- Generating raw waveforms at 16 kHz (very uncommon)

WaveNet: dilated ConvNet



[van der Oord et al, 2016]

WaveNet: dilated ConvNet



- Repeated pattern of dilations:

$1, 2, \dots, 512, 1, 2, \dots, 512, \dots$


- Gated (bilinear) non-linearity:

$$\mathbf{z} = \tanh(W_{f,k} * \mathbf{x}) \odot \sigma(W_{g,k} * \mathbf{x})$$

- There are also skip connections

[van der Oord et al, 2016]

Details of the ConvBlock



[van der Oord et al, 2016]

WaveNet: speech results

- Trained on 24.6 hours of speech
- Receptive field is 0.24 seconds
- Conditioned on the speaker ID



[van der Oord et al, 2016]

WaveNet: piano results

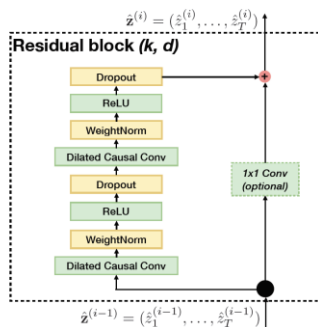



- Trained on 60 hours of piano (from YouTube)



[van der Oord et al, 2016]

Similar ConvNet for sequence modeling



[Bai et al. 2018]

ConvNets vs RNNs

Sequence Modeling Task	Model Size (\approx)	Models			
		LSTM	GRU	RNN	TCN
Seq. MNIST (accuracy ^{<i>h</i>})	70K	87.2	96.2	21.5	99.0
Permuted MNIST (accuracy)	70K	85.7	87.3	25.3	97.2
Adding problem $T=600$ (loss ^{ℓ})	70K	0.164	5.3e-5	0.177	5.8e-5
Copy memory $T=1000$ (loss)	16K	0.0204	0.0197	0.0202	3.5e-5
Music JSB Chorales (loss)	300K	8.45	8.43	8.91	8.10
Music Nottingham (loss)	1M	3.29	3.46	4.05	3.07
Word-level PTB (perplexity ^{ℓ})	13M	78.93	92.48	114.50	89.21
Word-level Wiki-103 (perplexity)	-	48.4	-	-	45.19
Word-level LAMBADA (perplexity)	-	4186	-	14725	1279
Char-level PTB (bpc ^{ℓ})	3M	1.41	1.42	1.52	1.35
Char-level text8 (bpc)	5M	1.52	1.56	1.69	1.45

[Bai et al. 2018]

Picking a probabilistic model

- N-grams
- CNNs (aka TDNNs)
- Any probabilistic classifier

Common problem: picking size of the window

- Avoiding overfitting
- To work on instances of different length
- To track long-range behavior

Probabilistic modeling of long sequences

$$p(x_M, x_{M-1}, \dots, x_1) =$$

$$p(x_M | x_{M-1}, \dots, x_1) \cdot p(x_{M-1}, x_{M-2}, \dots, x_1) =$$

$$\prod_{j=2}^M p(x_j | x_{j-1}, \dots, x_1) \cdot p(x_1) \approx$$

$$\prod_{j=1}^M p_{\theta}(x_j | h_{j-1})$$


“context variable”

$$h_{j-1} = F(x_{j-1}, h_{j-2})$$

Let us use a simple network here!

Recurrent neural network (RNN)

$$p_{\theta}(x_t^i \mid x_{t-1}^i, x_{t-2}^i, \dots, x_1^i)$$




$$h_t = W \phi(h_{t-1}) + W_x x_t$$

$$y_t = W_y \phi(h_t)$$

NB: I omit bias terms but they can be useful!

Most popular non-linearity for RNNs


$$\tanh x : \mathbb{R} \rightarrow (-1; 1)$$



$$[\tanh x]' = 1 - \tanh^2 x$$

$$|[\tanh x]'| \leq 1$$


Unwrapping RNN



$$h_t = W \phi(h_{t-1}) + W_x x_t$$

$$y_t = W_y \phi(h_t)$$

Training RNN



$$h_t = W \phi(h_{t-1}) + W_x x_t \quad y_t = W_y \phi(h_t)$$

Training RNN


$$h_t = W \phi(h_{t-1}) + W_x x_t$$

$$y_t = W_y \phi(h_t)$$

$$E = \sum_{t=1}^S E_t \quad \frac{dE}{dW} = \sum_{t=1}^S \frac{dE_t}{dW}$$


$$\frac{\partial E_t}{\partial W} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}$$

Training RNN



$$\frac{\partial E_t}{\partial W} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}$$

Training RNN



In practice: unwrapping for a finite number of time-steps (or training on bounded length sequences)

Training RNN

$$\frac{\partial E_t}{\partial W} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}$$

$$\frac{\partial h_t}{\partial h_k} = \prod_{i=k+1}^t \frac{\partial h_i}{\partial h_{i-1}} = \prod_{i=k+1}^t W^T \text{diag}(\phi'(h_{i-1}))$$

$$\left\| \frac{\partial h_i}{\partial h_{i-1}} \right\|_2 \leq \|W\|_2 L_\phi = \sigma_{\max} L_\phi$$

$$\left\| \frac{\partial h_i}{\partial h_k} \right\|_2 \leq (\sigma_{\max} L_\phi)^{t-k}$$

Challenges with training RNN

$$\frac{\partial E_t}{\partial W} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}$$

$$\left\| \frac{\partial h_i}{\partial h_k} \right\|_2 \leq (\sigma_{\max} L_\phi)^{t-k}$$

$$\sigma_{\max} L_\phi < 1$$

vanishing gradient:
network ignores
long links


$$\sigma_{\max} L_\phi > 1$$

exploding gradient:
learning quickly
“explodes”

Gradient clipping

Algorithm 1 Pseudo-code for norm clipping

```
 $\hat{\mathbf{g}} \leftarrow \frac{\partial \mathcal{E}}{\partial \theta}$   
if  $\|\hat{\mathbf{g}}\| \geq \text{threshold}$  then  
     $\hat{\mathbf{g}} \leftarrow \frac{\text{threshold}}{\|\hat{\mathbf{g}}\|} \hat{\mathbf{g}}$   
end if
```



- Simple trick handles gradient explosion (provided that the “valley” is wide)

[Pascanu et al. 2013]

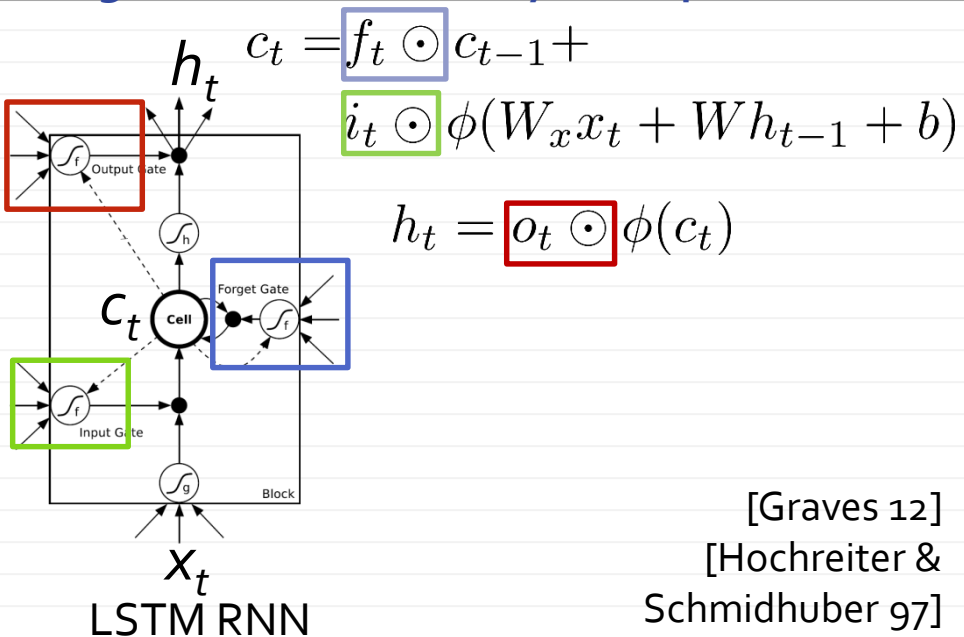
Handling vanishing gradient

$$\frac{\partial E_t}{\partial W} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}$$

$$\sigma_{\max} L_{\phi} < 1$$


- Even if the gradient does not vanish totally, the information stored in low-energy subspaces will not be propagated
- Idea: we need mechanism to ensure long-term propagation.

Long Short-term Memory: cell update



[Graves 12]
[Hochreiter &
Schmidhuber 97]

Long Short-term Memory: gate activations



$$o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1} + b_o)$$


$$f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1} + b_f)$$

$$i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1} + b_i)$$

LSTM RNN


[Graves 12]
[Hochreiter &
Schmidhuber 97]

Vanishing gradient visualization



The influence of an input unit quickly
vanishes with time [Graves 12]

Long Short-term Memory



○ = gate open

— = gate closed

[Graves 12]

[Hochreiter & Schmidhuber 97]

Recap: RNN-LSTM as a probabilistic model

$$p(x_t | x_1 \dots x_{t-1}) = ?$$

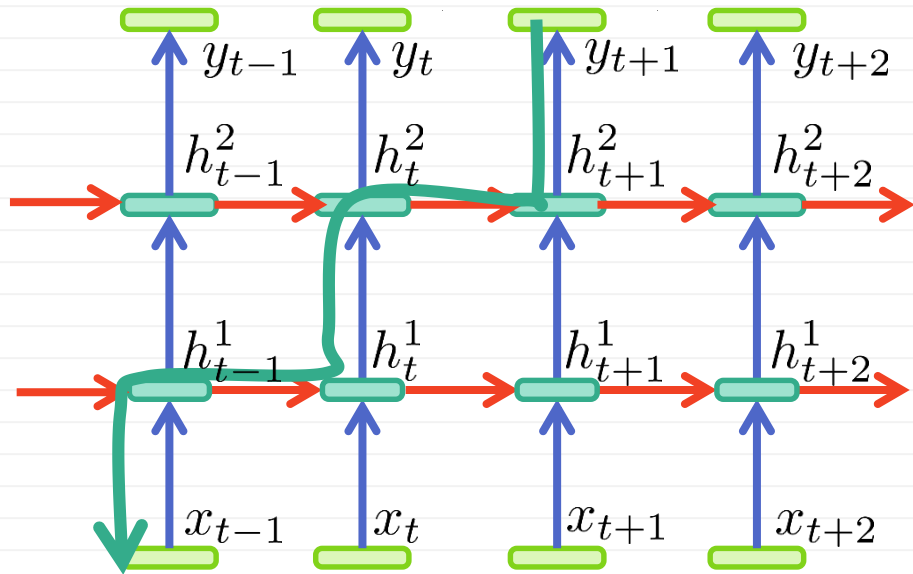
$$h_t = \text{LSTM}(x_{t-1}, h_{t-1})$$

$$y_t = W_y h_t$$

$$p_t^i = \frac{\exp(y_t^i)}{\sum_k \exp(y_t^k)} = p(i | x_1, \dots, x_{t-1})$$

$$x_t \sim \{p_t^i\}$$

Multi-layer RNNs



Computer generated “Linux kernel code”

© Andrej
Karpathy:

```
static void do_command(struct seq_file *m, void *v)
{
    int column = 32 << (cmd[2] & 0x80);
    if (state)
        cmd = (int)(int_state ^ (in_8(&ch->ch_flags) & Cmd) ? 2 : 1);
    else
        seq = 1;
    for (i = 0; i < 16; i++) {
        if (k & (1 << 1))
            pipe = (in_use & UMTHREAD_UNCCA) +
                ((count & 0x00000000ffffff8) & 0x000000f) << 8;
        if (count == 0)
            sub(pid, ppc_md.kexec_handle, 0x20000000);
        pipe_set_bytes(i, 0);
    }
    /* Free our user pages pointer to place camera if all dash */
    subsystem_info = &of_changes[PAGE_SIZE];
    rek_controls(offset, idx, &soffset);
    /* Now we want to deliberately put it to device */
    control_check_polarity(&context, val, 0);
    for (i = 0; i < COUNTER; i++)
        seq_puts(s, "policy ");
}
```

More fun: 1) <http://karpathy.github.io/2015/05/21/rnn-effectiveness/> 2) your assignment

Interpretable LSTM Cells

Cell sensitive to position in line:

The sole importance of the crossing of the Berezina lies in the fact that it plainly and indubitably proved the fallacy of all the plans for cutting off the enemy's retreat and the soundness of the only possible line of action--the one Kutuzov and the general mass of the army demanded--namely, simply to follow the enemy up. The French crowd fled at a continually increasing speed and all its energy was directed to reaching its goal. It fled like a wounded animal and it was impossible to block its path. This was shown not so much by the arrangements it made for crossing as by what took place at the bridges. When the bridges broke down, unarmed soldiers, people from Moscow and women with children who were with the French transport, all--carried on by vis inertiae--pressed forward into boats and into the ice-covered water and did not, surrender..

Cell that turns on inside quotes:

"You mean to imply that I have nothing to eat out of.... On the contrary, I can supply you with everything even if you want to give dinner parties," warmly replied Chichagov, who tried by every word he spoke to prove his own rectitude and therefore imagined Kutuzov to be animated by the same desire.

Kutuzov, shrugging his shoulders, replied with his subtle penetrating smile: "I meant merely to say what I said."

Cell that robustly activates inside if statements:

```
static int __dequeue_signal(struct sigpending *pending, sigset_t *mask,
                           siginfo_t *info)
{
    int sig = next_signal(pending, mask);
    if (sig) {
        if (current->notifier) {
            if (sigismember(current->notifier_mask, sig)) {
                if (!current->notifier(current->notifier_data)) {
                    clear_thread_flag(TIF_SIGPENDING);
                    return 0;
                }
            }
        }
        collect_signal(sig, pending, info);
    }
    return sig;
}
```

[Karpathy et al.
ICLR16]

Interpretable LSTM Cells

Cell that turns on inside comments and quotes:

```
/* duplicate LSM field information. The lsm_rule is opaque, so
 * re-initialized. */
static inline int audit_dupe_lsm_field(struct audit_field *df,
                                     struct audit_field *sf)
{
    int ret = 0;
    char *lsm_str;
    /* our own copy of lsm_str */
    lsm_str = kstrdup(sf->lsm_str, GFP_KERNEL);
    if (unlikely(!lsm_str))
        return -ENOMEM;
    df->lsm_str = lsm_str;
    /* our own (refreshed) copy of lsm_rule */
    ret = security_audit_rule_init(df->type, df->op, df->lsm_str,
                                  (void **)&df->lsm_rule);
    /* keep currently invalid fields around in case they
     * become valid after a policy reload. */
    if (ret == -EINVAL) {
        pr_warn("audit rule for LSM '%s' is invalid\n",
                df->lsm_str);
        ret = 0;
    }
    return ret;
}
```

[Karpathy et al.
ICLR16]

Interpretable LSTM Cells

Cell that is sensitive to the depth of an expression:

```
#ifdef CONFIG_AUDITSYSCALL
static inline int audit_match_class_bits(int class, u32 *mask)
{
    int i;
    if (classes[class]) {
        for (i = 0; i < AUDIT_BITMASK_SIZE; i++)
            if (mask[i] & classes[class][i])
                return 0;
    }
    return 1;
}
```

Cell that might be helpful in predicting a new line. Note that it only turns on for some “)”:

```
char *audit_unpack_string(void **bufp, size_t *remain, si
{
    char *str;
    if (!*bufp || (len == 0) || (len > *remain))
        return ERR_PTR(-EINVAL);
    /* Of the currently implemented string fields, PATH_MAX
     * defines the longest valid length.
     */
    if (len > PATH_MAX)
        return ERR_PTR(-ENAMETOOLONG);
    str = kmalloc(len + 1, GFP_KERNEL);
    if (unlikely(!str))
        return ERR_PTR(-ENOMEM);
    memcpy(str, *bufp, len);
    str[len] = 0;
    *bufp += len;
    *remain -= len;
    return str;
}
```

[Karpathy et al.
ICLR16]

Non-interpretable LSTM Cells


A large portion of cells are not easily interpretable. Here is a typical example:

```
/* Unpack a filter field's string representation from user-space
 * buffer. */
char *audit_unpack_string(void **bufp, size_t *remain, size_t len)
{
    char *str;
    if (!*bufp || (len == 0) || (len > *remain))
        return ERR_PTR(-EINVAL);
    /* of the currently implemented string fields, PATH_MAX
     * defines the longest valid length.
     */
}
```

[Karpathy et al.
ICLR16]

Interpretable LSTM Cells


<0.1
(do not copy)




[Karpathy et al.
ICLR16]

>0.9 (copy intact)

Gated Recurrent Units (GRU)



(a) Long Short-Term Memory



(b) Gated Recurrent Unit

[Cho et al. 14]

[Chung et al. 14]

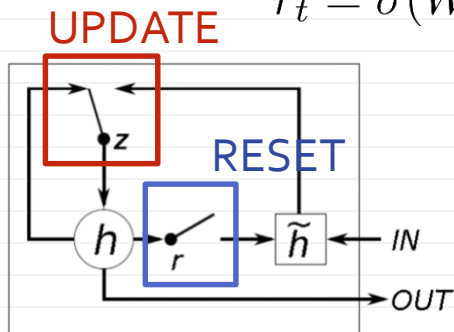
Gated Recurrent Units (GRU)

$$h_t = (1 - z_t) \odot h_{t-1} +$$

$$z_t \odot \phi(W_x x_t + W_r r_t \odot h_{t-1} + b)$$

$$z_t = \sigma(W_{xz}x_t + W_{hz}h_{t-1} + b_z)$$


$$r_t = \sigma(W_{xr}x_t + W_{hr}h_{t-1} + b_r)$$



[Cho et al. 14]

[Chung et al. 14]

GRU gate statistics



[Karpathy et al. ICLR16]

Plain vs LSTM vs GRU

Success at predicting next characters in the test sequence (cross-entropy loss) [Karpathy et al. ICLR16]:


	LSTM			RNN			GRU		
Layers	1	2	3	1	2	3	1	2	3
Size	War and Peace Dataset								
64	1.449	1.442	1.540	1.446	1.401	1.396	1.398	1.373	1.472
128	1.277	1.227	1.279	1.417	1.286	1.277	1.230	1.226	1.253
256	1.189	1.137	1.141	1.342	1.256	1.239	1.198	1.164	1.138
512	1.161	1.092	1.082	-	-	-	1.170	1.201	1.077
	Linux Kernel Dataset								
64	1.355	1.331	1.366	1.407	1.371	1.383	1.335	1.298	1.357
128	1.149	1.128	1.177	1.241	1.120	1.220	1.154	1.125	1.150
256	1.026	0.972	0.998	1.171	1.116	1.116	1.039	0.991	1.026
512	0.952	0.840	0.846	-	-	-	0.943	0.861	0.829

Model \ n	1	2	3	4	5	6	7	8	9	20
	War and Peace Dataset									
n -gram	2.399	1.928	1.521	1.314	1.232	1.203	1.194	1.194	1.194	1.195
n -NN	2.399	1.931	1.553	1.451	1.339	1.321	-	-	-	-
	Linux Kernel Dataset									
n -gram	2.702	1.954	1.440	1.213	1.097	1.027	0.982	0.953	0.933	0.889
n -NN	2.707	1.974	1.505	1.395	1.256	1.376	-	-	-	-


RNN success cases

Success at predicting next characters in the test sequence (cross-entropy loss) [Karpathy et al. ICLR16]:

Linux kernel:




War and peace:



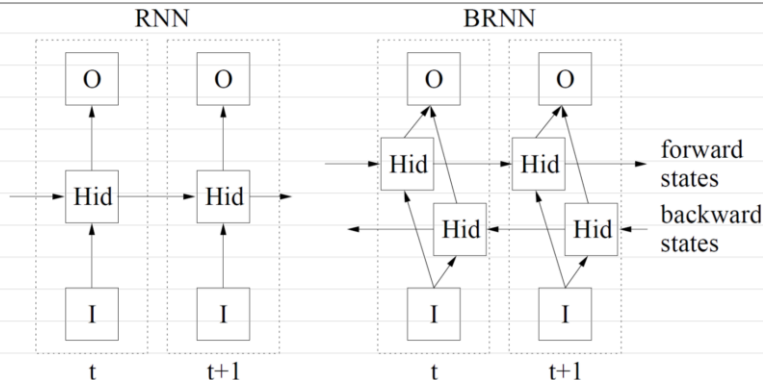
Plain vs LSTM vs GRU

Model similarity (t-SNE embedding of character probabilities):



[Karpathy et al. ICLR16]

Bi-directional RNN



for $t = 1$ to T do

Do forward pass for the forward hidden layer, storing activations at each timestep

for $t = T$ to 1 do

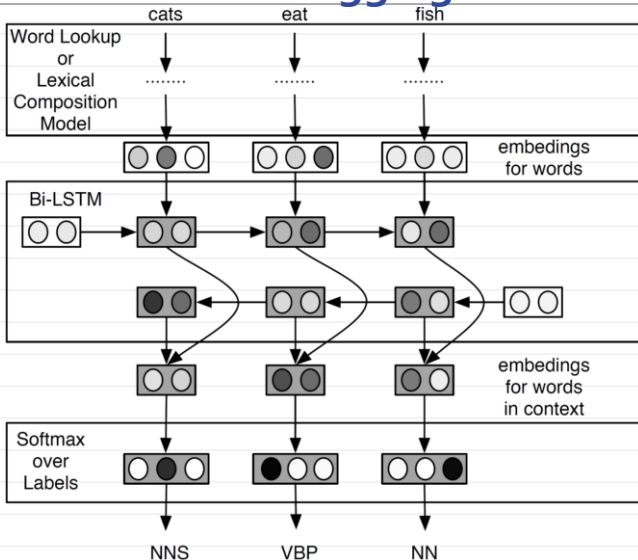
Do forward pass for the backward hidden layer, storing activations at each timestep

for $t = 1$ to T do

Do forward pass for the output layer, using the stored activations from both hidden layers

[A Graves, PhD thesis]

Bi-LSTM POS tagging



$$\mathbf{l}_i = \tanh(\mathbf{L}^f \mathbf{s}_i^f + \mathbf{L}^b \mathbf{s}_i^b + \mathbf{b}_l)$$

[Ling et al.
EMNLP₁₅]

Bi-LSTM POS tagging


	acc	parameters	words/sec
Word Lookup	96.97	2000k	6K
Convolutional (S&Z)	96.80	42.5k	4K
Forward RNN	95.66	17.5k	4K
Backward RNN	95.52	17.5k	4K
Bi-RNN	95.93	40k	3K
Forward LSTM	97.12	80k	3K
Backward LSTM	97.08	80k	3K
Bi-LSTM $d_{CS} = 50$	97.22	70k	3K
Bi-LSTM	97.36	150k	2K

[Ling et al. EMNLP15]

Uni-directional vs bi-directional

- Bi-directional is not applicable when “future” is unavailable
- When future is available bi-directional is almost always better
- E.g. NLP (batch mode), bioinformatics

BiLM/ELMo



- BiLM is learned for language modeling (next word prediction)
- ELMo coefficients are fine-tuned for the new task:

$$\text{ELMo}_k^{\text{task}} = E(R_k; \Theta^{\text{task}}) = \gamma^{\text{task}} \sum_{j=0}^L s_j^{\text{task}} \mathbf{h}_{k,j}^{\text{LM}}$$

[Peters et al. ACL18]

BiLM/ELMO

Source		Nearest Neighbors
GloVe	play	playing, game, games, played, players, plays, player, Play, football, multiplayer
biLM	Chico Ruiz made a spectacular <u>play</u> on Alusik 's grounder { ... }	Kieffer , the only junior in the group , was commended for his ability to hit in the clutch , as well as his all-round excellent <u>play</u> .
	Olivia De Havilland signed to do a Broadway <u>play</u> for Garson { ... }	{ ... } they were actors who had been handed fat roles in a successful <u>play</u> , and had talent enough to fill the roles competently , with nice understatement .

Table 4: Nearest neighbors to “play” using GloVe and the context embeddings from a biLM.

SNLI task [Bowman et al. EMNLP15]:

Text	Judgments	Hypothesis
A man inspects the uniform of a figure in some East Asian country.	contradiction C C C C C	The man is sleeping
An older and younger man smiling.	neutral N N E N N	Two men are smiling and laughing at the cats playing on the floor.
A black race car starts up in front of a crowd of people.	contradiction C C C C C	A man is driving down a lonely road.
A soccer game with multiple males playing.	entailment E E E E E	Some men are playing a sport.
A smiling costumed woman is holding an umbrella.	neutral N N E C N	A happy woman in a fairy costume holds an umbrella.




Figure 1: Comparison of baseline vs. ELMo performance for SNLI and SRL as the training set size is varied from 0.1% to 100%.

[Peters et al. ACL18]

Recap and outlook

- Sequence prediction
- ConvNets and Recurrent nets are SOA in wide variety of domains (NLP, speech/signal, bioinformatics)
- Gating in RNN makes its memory longer
- Sequence prediction extends to other tasks (fixed->seq, seq->fixed, seq->seq, fixed->fixed) – *next lecture*

Bibliography

Elman, Jeffrey L. Finding structure in time. Cognitive science, 14(2):179–211, 1990.

A. Karpathy

The Unreasonable Effectiveness of Recurrent Neural Networks

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Razvan Pascanu, Tomas Mikolov, Yoshua Bengio:

On the difficulty of training recurrent neural networks. ICML (3) 2013: 1310-1318

A. Graves. Supervised Sequence Labelling with Recurrent Neural Networks. Textbook, Studies in Computational Intelligence, Springer, 2012

Sepp Hochreiter, Jürgen Schmidhuber:

Long Short-Term Memory. Neural Computation 9(8): 1735-1780 (1997)

Bibliography

Kyunghyun Cho, Bart van Merriënboer, Çağlar Gülçehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, Yoshua Bengio: Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation. EMNLP 2014: 1724-1734

Junyoung Chung, Çağlar Gülçehre, KyungHyun Cho, Yoshua Bengio: Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling. CoRRabs/1412.3555 (2014)

Andrej Karpathy, Justin Johnson, Fei-Fei Li: Visualizing and Understanding Recurrent Networks. ICLR 2016

Mikolov, T., Joulin, A., Chopra, S., Mathieu, M., & Ranzato, M. A. (2014). Learning longer memory in recurrent neural networks. arXiv preprint arXiv:1412.7753.

Aäron van den Oord, Sander Dieleman, Heiga Zen, Karen Simonyan, Oriol Vinyals, Alex Graves, Nal Kalchbrenner, Andrew W. Senior, Koray Kavukcuoglu: WaveNet: A Generative Model for Raw Audio. SSW 2016: 125

Bibliography

Wang Ling, Chris Dyer, Alan W. Black, Isabel Trancoso, Ramon Fernandez, Silvio Amir, Luís Marujo, Tiago Luís: Finding Function in Form: Compositional Character Models for Open Vocabulary Word Representation. EMNLP 2015: 1520-1530

Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, Luke Zettlemoyer:
Deep Contextualized Word Representations. NAACL-HLT 2018: 2227-2237